



Melbourne Bioinformatics

BIOINFORMATICS + DATA SERVICES + INFRASTRUCTURE, FOR LIFE SCIENCES TODAY



COMP90014

Algorithms for Bioinformatics

Week 10B: Unsupervised Learning - Clustering II

Machine learning: clustering

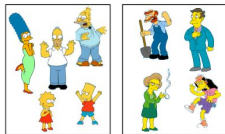
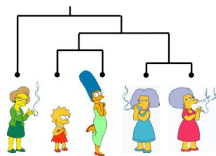
1. Clustering approaches
2. Density-based clustering
3. Divisive clustering

Clustering approaches

Exclusive Clustering



Overlapping Clustering



Exclusive

- ☑ Data points belong to only one cluster

Overlapping

- ☑ Data points may belong to many clusters

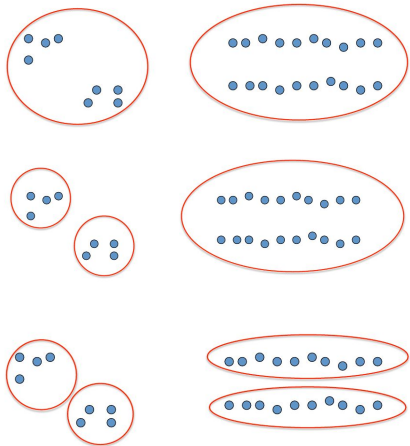
Hierarchical

- ☑ Assign points to “nested” clusters
- ☑ Get all possible clusters for given metric

Partitional

- ☑ Split points into “flat” independent clusters
- ☑ How many?

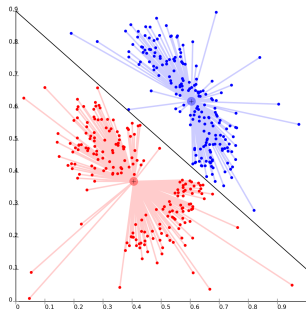
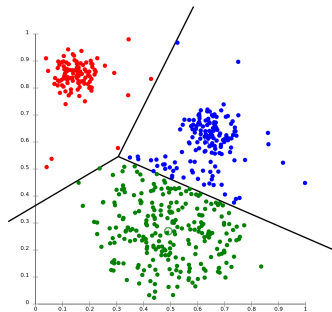
Ingredients for clustering



1. similarity metric
 - e.g. Euclidean distance
 2. a function to evaluate the quality of the clusters
 3. clustering algorithm
- 🧠 clustering is subjective
e.g. how many clusters?
- fixed k clusters
 - find the best k to optimize a function

Clustering approaches by cluster definition


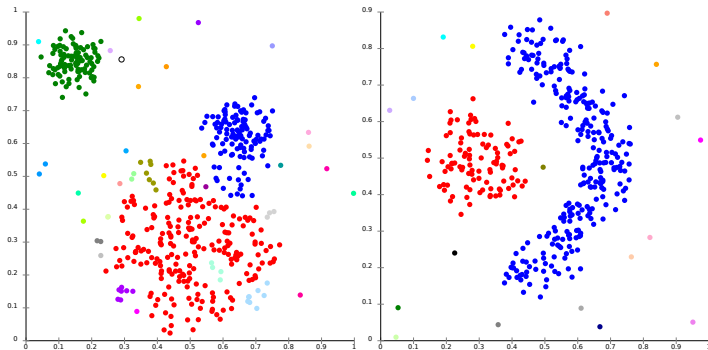
1. Centroid-based (k -means, k -medoids)



Notion of clusters: Voronoi tessellation/diagram

Clustering approaches by cluster definition



2. Connectivity-based (hierarchical)



Notion of clusters: cut dendrogram at some depth

Clustering approaches by cluster definition

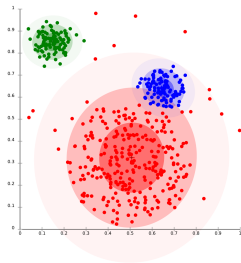
3. Density-based (DBSCAN, OPTICS)



Notion of clusters: connected regions of high density


Clustering approaches by cluster definition

4. Distribution-based (Mixture Models)




Notion of clusters: distributions on features

5. Network-based




Notion of clusters: graph connectivity

When k -means fails




(a) Original points.


(b) Three K-means clusters.



(a) Original points.




(b) Three K-means clusters.




(A): Undesirable clusters


(B): Ideal clusters



(a) Original points.



(b) Two K-means clusters.




(A): Two natural clusters

(B): k -means clusters

Machine learning: clustering

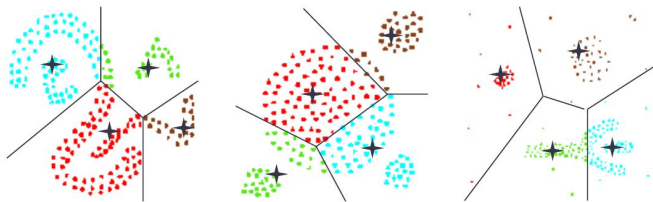
1. Clustering approaches
2. Density-based clustering
3. Divisive clustering

Concepts



- clusters are ***dense regions in the data space***, separated by regions of lower object density
- for any point in a cluster, the local point density around that point has to exceed some threshold (ϵ)
- the set of points from one cluster is spatially connected

Density-based clustering

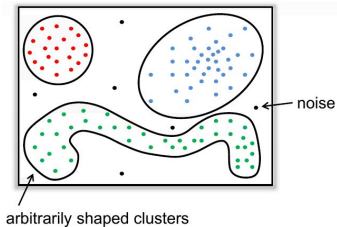



☹ When k -means fails ($K = 4$)

☹ DBSCAN:

- discover clusters of arbitrary shape
- handle noise and outliers

DBSCAN



- first density-based clustering algorithm
- one of the most widely used/cited clustering algorithms

Intuition:

- a cluster is a region of high density**
- noise points lie in regions of low density

We need to:

- define neighbourhood of a data point
- define high density

Definitions

ϵ -neighbourhood: objects within a radius ϵ of an object.


ϵ : **input parameter.**

High-density: ϵ -neighbourhood of an object contains at least minpts of objects.

minpts: **input parameter.**


$$N_{\epsilon}(p) : \{q \mid d(p, q) \leq \epsilon\}$$

ϵ -neighbourhood of p and q : Density of p is “high” (minpts = 4);
Density of q is “low” (minpts = 4)



Definitions

Given ϵ and minpts, categorize the objects into three exclusive categories.




Core point: It has more than minpts objects within ϵ .

Border point: It has fewer than minpts within ϵ , but is in the neighbourhood of a core point.

Noise/outlier point: Any point that is not a core point nor a border point.

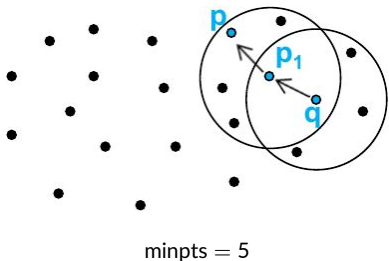
Direct density reachability



An object p is **directly density-reachable** from object q if:

- q is a core object and
- p is in its ϵ -neighborhood
- is p directly density-reachable from q ?
- is q directly density-reachable from p ?
- Density-reachability is an asymmetric relationship

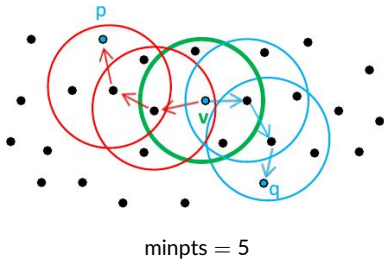
Density reachability



A point p is **density-reachable** from a point q if:

- There is a chain of points p_1, p_2, \dots, p_k , with $p_1 = q$ and $p_k = p$, such that p_{i+1} is directly density-reachable for all $1 < i < k - 1$
- Asymmetric

Density connectivity




A point p is *density-connected* to a point q if:

- 📍 there is a point v , such that both p and q are *density-reachable* from v
- 📍 Symmetric

Cluster definition

Given a data set D of points, parameter ϵ and minpts:

- a cluster C is a subset of D satisfying two criteria.



Maximality:

- $\forall p, q$ if $p \in C$
and if q is density-reachable from p ,
then also $q \in C$

Connectivity:

- $\forall p, q \in C$,
 p and q are density-connected

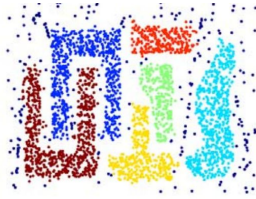
- clusters will contain both core and border points
- noise/outliers:
 - points in D not belonging to any cluster

DBSCAN algorithm

```
1 Procedure Dbscan( $X, \epsilon, \text{minpts}$ ):
2   foreach unvisited point  $x \in X$  do
3     mark  $x$  as visited
4      $N \leftarrow \text{GetNeighbours}(x, \epsilon)$ 
5     if  $|N| < \text{minpts}$  then
6       mark  $x$  as noise
7     else
8        $C \leftarrow \{x\}$ 
9       foreach point  $x' \in N$  do
10         $N \leftarrow N \setminus x'$ 
11        if  $x'$  is not visited then
12          mark  $x'$  as visited
13           $N' \leftarrow \text{GetNeighbours}(x', \epsilon)$ 
14          if  $|N'| \geq \text{minpts}$  then
15             $N \leftarrow N \cup N'$ 
16          if  $x'$  is not yet member of any
17            cluster then
               $C \leftarrow C \cup \{x'\}$ 
```

- ☞ Input parameters:
 - X points, ϵ and minpts
- ☞ the algorithm proceeds by arbitrarily picking (scanning) up points in the dataset until all points have been visited
- ☞ if p is a core point (at least minpts points within a radius of ϵ) collect all density-reachable points from p and assign to a new cluster
- ☞ assign p to noise otherwise
- ☞ don't change p 's cluster assignment

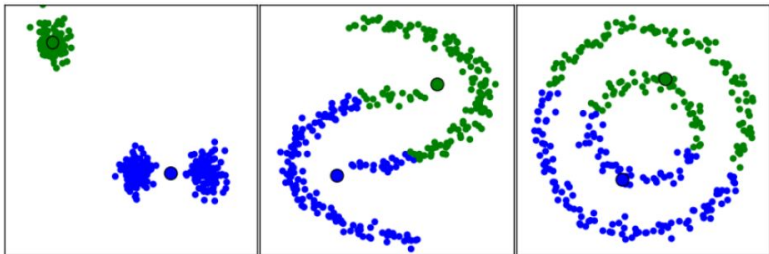
Complexity and strength of DBSCAN



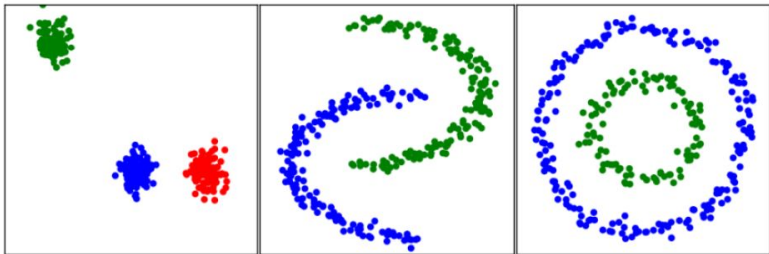
- ⌚ time complexity:
 - $O(n^2)$ if done naïvely
 - $O(n \times \log n)$ with a spatial index
 - only works in relatively low dimensions
- ⌚ space complexity: $O(n)$
- ⌚ can handle arbitrary shapes
- ⌚ can handle clusters of different sizes
- ⌚ resistant to noise

DBSCAN vs. *k*-means




K-means



DBSCAN



Weaknesses of DBSCAN



Goal:

- 🧠 varying densities
- 🧠 high dimensional data
- 🧠 overlapping clusters

Different ϵ configurations:

- 🧠 setting ϵ and minpts can be tricky

Determining ϵ





Figure 2 Points sorted by distance to the 3rd nearest neighbor



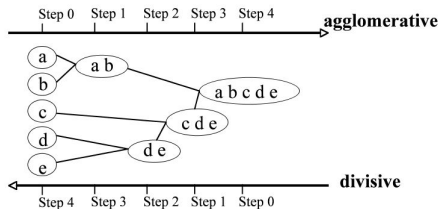
k-distance: calculate distance of k^{th} nearest neighbor for each point
 $k = \text{minpts} - 1$

- plot in ascending / descending order
- set ϵ to the maximum distance before the threshold
- noise points have their k^{th} nearest neighbour at higher distances

Machine learning: clustering

1. Clustering approaches
2. Density-based clustering
3. Divisive clustering

Hierarchical clustering



Agglomerative clustering (bottom-up)

- ☛ each data point starts as a single cluster
- ☛ join clusters into bigger clusters till we reach one single cluster with all points

Divisive clustering (top-down)

- ☛ start with one big cluster
- ☛ at each step, split into smaller clusters
- ☛ stop at desired number of clusters
- ☛ e.g. when points are in single clusters


What are the next clusters to merge?

Single linkage: $D_{k,g} = \min(D_{k,i}, D_{k,j})$

Complete linkage: $D_{k,g} = \max(D_{k,i}, D_{k,j})$

Average linkage: $D_{k,g} = \frac{D_{k,i} + D_{k,j}}{2}$

Divisive hierarchical clustering



- Any partitional algorithm that generates a fixed number of clusters can be used to implement divisive hierarchical clustering
 - e.g. k -means, with $k = 2$
 - keep partitioning clusters iteratively

Challenge: use k -means to implement divisive hierarchical clustering on a set of points X .

You can use assume the function `kmeans()` is available (you don't need to implement it yourself).

Hint: start by dividing X into two clusters, then recursively run `kmeans()` on the output until each cluster has only 1 item.

Thank you!

Today: Unsupervised Learning - Clustering II

Next time: Supervised Learning